

Advanced OpenMP

Advanced clauses, variable types, race conditions, task parallelism

OpenMP variable types

- When threads are created, an existing variable **x** in the master thread can have different status in the newly created threads, among which:
- **private(x)**: create private copy in each thread, uninitialized
- **firstprivate(x)**: create private copy in each thread, copy initial value
- **threadprivate(x)**: as if having a **global** copy of a variable per thread
- **shared(x)**: x of the master is shared among all threads for read/write.
- **default(shared | none)**: Variables **shared** by default, **none** disables it.
 - Good practice for preventing careless bugs
- **reduction(op:x)**: Creates a safe copy for each thread, merges them transparently at the end. Avoids race conditions.

Example: OpenMP variable types

OpenMP loop scheduling

- Specified with the clause **schedule**(schedule, chunksize)
- **static** schedule
 - Gives chunksize contiguous iterations to threads in a round-robin manner (th0, th1, ..., thP-1, th0, th1, ...)
 - If no chunksize specified, it is N/P by default (single chunk per thread)
 - If no schedule is specified, it is static with chunksize= N/P by default

OpenMP loop scheduling

- Specified with the clause **schedule**(schedule, chunksize)
- **dynamic** schedule
 - Gives chunksize contiguous iterations to the first available thread (th0, th1, th0, th3, th1, ...)
 - If no chunksize specified, it is 1 by default (single iteration per chunk)
 - Good to make sure that no thread is idling
 - Default chunksize is 1
 - Bad for data locality per thread

OpenMP loop scheduling

- Specified with the clause **schedule**(schedule, chunksize)
- **guided** schedule
 - Chunks are distributed dynamically to threads just like **dynamic** (th0, th1, th0, th3, th1, ...)
 - Starts with big chunks, but progressively reduces the size to the given chunksize towards the end
 - Makes sure all threads have some workload on the entire iteration domain
 - Keeps good locality for the most part
 - Compromise between **static** and **dynamic**

Example: Loop scheduling

Merging **parallel** and **for/sections** directives

- If the parallel region has an **omp for** or **omp sections** inside, and **nothing else**, you can merge two directives.
- You cannot use curly braces ({ ... }) after **#pragma omp parallel for**.

```
#pragma omp parallel num_threads(P)
{
    #pragma omp for
    for (int i = 0; i < N; i++) {
        f(i)
    }
}
```

```
#pragma omp parallel for num_threads(P)
for (int i = 0; i < N; i++) {
    f(i)
}
```

```
#pragma omp parallel sections num_threads(P)
{
    #pragma omp section
    {
        g();
    }
    #pragma omp section
    {
        h();
    }
}
```


collapse clause for omp for

- **collapse** clause flattens nested loops into a single domain to provide more parallelism (e.g., when iteration domain is small).
- Iteration domain must be a fixed size for all nested loops.
- Can collapse multiple loops (≥ 2)

```
#pragma omp parallel num_threads(P)
{
    // This loop is
    #pragma omp for collapse(2)
    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            f(i, j);
        }
    }
}
```

```
// Same as this
#pragma omp for
for (int i = 0; i < N*N; i++) {
    f(i/N, i%N);
}
}
```

ordered clause for omp for

- Specifies a block in a parallel loop that respects the sequential order of iterations in execution.
- Useful for gathering results in order at the end, after performing expensive computations in parallel.

```
// We want to have result[i] = f(i) in the end  
std::vector<int> result;
```

```
#pragma omp parallel num_threads(P)  
{  
  #pragma omp for ordered  
    for (int i = 0; i < N; i++) {  
      int res = f(i); // Done in parallel, time consuming  
    }  
  #pragma omp ordered  
    result.push_back(res); // Done in order sequentially  
}
```

Race conditions

- What is the final value of n?
min/max?
- There are three hidden instructions:
 - **load** n
 - **add** n, 1
 - **store** n
- Threads simultaneously read the old value before updating n
 - Some increments lost
 - => **race condition**
- Solutions?
 - **atomic**
 - **critical**
 - **reduction**

```
// We want to have result[i] = f(i) in the end  
int n = 0;
```

```
#pragma omp parallel for num_threads(P)  
  for (int i = 0; i < 1000000; i++) {  
    n++;  
  }
```

Race conditions

- What is the final value of n? min/max?
- There are three hidden instructions:
 - **load** n
 - **add** n, 1
 - **store** n
- Threads simultaneously read the old value before updating n
 - Some increments lost
 - => **race condition**
- Solutions?
 - **atomic** directive
 - Uses hardware atomic instructions for the following instruction, has some overhead
 - Works for basic arithmetic / logical operations (+, -, *, /, min, max, &, |, ...)

```
// We want to have result[i] = f(i) in the end  
int n = 0;
```

```
#pragma omp parallel for num_threads(P)  
    for (int i = 0; i < 1000000; i++) {  
#pragma omp atomic  
        n++;  
    }
```

Race conditions

- What is the final value of n? min/max?
- There are three hidden instructions:
 - **load** n
 - **add** n, 1
 - **store** n
- Threads simultaneously read the old value before updating n
 - Some increments lost
 - => **race condition**
- Solutions?
 - **critical** directive
 - Block of code in a **critical** region is executed by a **single thread at a time**
 - Works for all code types but higher overhead than **atomic**

```
// We want to have result[i] = f(i) in the end  
int n = 0;
```

```
#pragma omp parallel for num_threads(P)  
    for (int i = 0; i < 1000000; i++) {  
#pragma omp critical  
    {  
        n++;  
    }  
}
```

Race conditions

- What is the final value of n ? min/max?
- There are three hidden instructions:
 - **load** n
 - **add** $n, 1$
 - **store** n
- Threads simultaneously read the old value before updating n
 - Some increments lost
 - => **race condition**
- Solutions?
 - **reduction(op:x)** clause
 - Creates an invisible copy of the shared variable x in each thread, performs the reduction operation **op** on the private copy without atomic nor critical.
 - At the end of the loop, merges these private copies with a single atomic operation **op** on the shared x .
 - **op** can be $+, -, *, /, \min, \max, \&, |, \&\&, | |, ^, \dots$.

```
// We want to have result[i] = f(i) in the end  
int n = 0;
```


```
#pragma omp parallel for num_threads(P) reduction(+:n)  
    for (int i = 0; i < 1000000; i++) {  
        n++;  
    }  
}
```

OpenMP task parallelism

Task creation, management, and dependencies

Task-based parallelism

- A task is a block of computation (block of code, function, even loop iterations)
- A task might need another task's output, which is called a **dependency**.
- Task-based parallelism expresses the computation as a bunch of tasks, and their interdependencies.
- A **runtime system** then orchestrates their execution among multiple threads.
- It is particularly useful for
 - irregular parallel computations
 - recursive functions
- Advanced task-based runtime systems enable hybrid GPU+CPU+FPGA parallelism
 - For each task, specify data dependencies
 - For each device/task, specify code/kernel
 - Runtime system orchestrates the distribution of tasks to compute units



Creating tasks in OpenMP

- When a thread encounters a **#pragma omp task** directive, it creates a task involving the following code block, and puts it into a task pool instead of executing it.
- When one of the threads in the parallel region is idle, it executes one of the **ready** tasks (whose all dependencies are satisfied).
- No particular order for task execution, aside from dependencies.
- The program on the right, is it correct?

```
#pragma omp parallel num_threads(P)
{
    #pragma omp task
    f();
    #pragma omp task
    g();
}
```

Creating tasks in OpenMP

- Each task must be created by a **single thread**
 - Must use a work sharing construct
- **#pragma omp taskwait** waits until all tasks created by current thread/task are done
 - It **does not** wait for all tasks globally.
- **#pragma omp barrier** implies a **taskwait**
- The program on the right, is it correct?

```
#pragma omp parallel num_threads(P)
{
    #pragma omp single
    {
        #pragma omp task
        f();
        #pragma omp task
        g();
        #pragma omp taskwait
        // After taskwait f() and g() are guaranteed to
        finish.
    }
    // Here, even without taskwait f() and g() are executed.
}

#pragma omp parallel for num_threads(P)
for (int i = 0; i < N; i++) {
    #pragma omp task
    h(i);
}
```

Task dependencies

- Tasks can have dependencies on variables.
- Dependency can be
 - **input** dependency: Needed to start the task
 - **depend(in:var)** clause after **task**
 - **output** dependency: Other tasks having an input dependency on this variable must wait for this task.
 - **depend(out:var)** clause after **task**
- It is the programmer's responsibility to correctly specify dependencies.

```
// We want to compute x = f() + g();
int x;
int y;

#pragma omp parallel num_threads(P)
{
    #pragma omp single
    {
        #pragma omp task depend(out:x)
        x = f();
        #pragma omp task depend(out:y)
        y = g();
        #pragma omp task depend(in:x) depend(in:y)
        x = x + y;
        #pragma omp taskwait
        printf("f() + g() = %d\n", x);
    }
}
```

Example: Computing nine